

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

In [2]: housedf=pd.read_csv("USA_Housing.csv")

In [3]: housedf.head()

Out[3]:   Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  Avg. Area Number of Bedrooms  Area Population  Price  Address
0      79545.458574         5.682861         7.009188              4.09      23086.800503  1.059034e+06  208 Michael Ferry Apt. 674\inLaurabury, NE 3701...
1      79248.642455         6.002900         6.730821              3.09      40173.072174  1.505891e+06  188 Johnson Views Suite 079\inLake Kathleen, CA...
2      61287.067179         5.865890         8.512727              5.13      36882.159400  1.059988e+06  9127 Elizabeth Stravenue\inDanielstown, WI 06482...
3      63345.240046         7.188236         5.586729              3.26      34310.242831  1.260617e+06              USS Barnett\inFPO AP 44820
4      59982.197226         5.040555         7.839388              4.23      26354.109472  6.309435e+05              USNS Raymond\inFPO AE 09386

In [4]: housedf.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Avg. Area Income     5000 non-null  float64
1   Avg. Area House Age  5000 non-null  float64
2   Avg. Area Number of Rooms  5000 non-null  float64
3   Avg. Area Number of Bedrooms  5000 non-null  float64
4   Area Population       5000 non-null  float64
5   Price                5000 non-null  float64
6   Address              5000 non-null  object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB

In [5]: housedf.describe()

Out[5]:   Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  Avg. Area Number of Bedrooms  Area Population  Price
count      5000.000000         5000.000000         5000.000000         5000.000000      5000.000000  5.000000e+03
mean      68583.108984         5.977222         6.987792              3.981330      36163.516039  1.232073e+06
std       10657.991214         0.991456         1.005833              1.234137      9925.650114  3.531176e+05
min       17796.631190         2.644304         3.236194              2.000000      172.610686  1.593866e+04
25%       61480.562388         5.322283         6.299250              3.140000      29403.928702  9.975771e+05
50%       68804.286404         5.970429         7.002902              4.050000      36199.406689  1.232669e+06
75%       75783.338666         6.650808         7.665871              4.490000      42861.290769  1.471210e+06
max       107701.748378         9.519088         10.759588              6.500000      69621.713378  2.469066e+06

In [6]: housedf.columns

Out[6]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
      'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')

In [7]: sns.pairplot(housedf)

Out[7]: <seaborn.axisgrid.PairGrid at 0x26f96e15b50>

In [8]: sns.heatmap(housedf.corr(), annot=True)

Out[8]: <AxesSubplot:~>

In [9]: X= housedf[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
      'Avg. Area Number of Bedrooms', 'Area Population']]
y=housedf['Price']

In [10]: from sklearn.model_selection import train_test_split

In [11]: X_train, X_test, y_train, y_test=train_test_split(X, y, train_size=0.40 , random_state=101)

In [12]: X_train

Out[12]:   Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  Avg. Area Number of Bedrooms  Area Population
1672      49775.405947         5.305108              6.178535              2.24      23557.361654
4557      75571.044023         6.114928              7.318214              4.44      33988.435859
1040      66250.316685         4.858520              7.758692              3.08      24201.753077
514       58978.222191         6.658346              7.889948              3.01      44071.604628
2148      73053.953218         5.712999              6.673142              4.30      41686.518927
...
4171      56610.642563         4.846832              7.558137              3.29      25494.740298
599       70596.850945         6.548274              6.539986              3.10      51614.830136
1361      55621.899104         3.735942              6.868291              2.30      63184.613147
1547      63044.460096         5.935261              5.913454              4.10      32725.279544
4959      75078.791516         7.644779              8.440726              4.33      56148.449322
2000 rows x 5 columns

In [13]: y_train

Out[13]: 1672      4.540557e+05
4557      1.218264e+06
1040      1.094322e+06
514       1.308984e+06
2148      1.220724e+06
...
4171      7.296417e+05
599       1.599479e+06
1361      1.192641e+06
1547      8.650995e+05
4959      2.108376e+06
Name: Price, Length: 2000, dtype: float64

In [14]: from sklearn.linear_model import LinearRegression

In [11]: Lr=LinearRegression()

In [12]: Lr.fit(X_train, y_train)

Out[12]: LinearRegression()

In [13]: coeff_df=pd.DataFrame(Lr.coef_, X.columns,columns=['coefficient'])

In [14]: coeff_df

Out[14]:   coefficient
Avg. Area Income     21.384444
Avg. Area House Age  162975.483408
Avg. Area Number of Rooms  121802.121132
Avg. Area Number of Bedrooms  1934.163056
Area Population      15.189607

In [15]: Prediction= Lr.predict(X_test)

In [31]: plt.scatter(y_test, Prediction)

Out[31]: <matplotlib.collections.PathCollection at 0x2635cb097f0>

In [38]: sns.distplot((y_test-Prediction),bins=70)

C:\Users\918130\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: 'distplot' is a deprecated function and will be removed in a future version. Please adapt y
our code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[38]: <AxesSubplot: xlabel='Price', ylabel='Density'>

In [34]: from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, Prediction)
r_square = r2_score(y_test, Prediction)

In [35]: print('Mean_Squared_Error :', mse)
print('r_square_value : ',r_square)

Mean_Squared_Error : 10354334181.172644
r_square_value : 0.9175452486480594

In [36]: from math import sqrt
rms = sqrt(mse)
rms

Out[36]: 101756.24885564839

In [ ]:
```